

Sachin Koli

Solution Architect — Cloud Data & AI Platforms

Google Cloud Certified — Professional Data Engineer

Pune, India · +91 95522 36200 · kolisachint@gmail.com · linkedin.com/in/kolisachint · github.com/kolisachint

WHAT I DO

I architect enterprise data platforms on Google Cloud. Right now that means moving a **UK tier-1 retail bank** off its enterprise data warehouse onto a GCP-native platform against a **2026 decommissioning date** — event-driven ingestion on **Confluent Kafka** and **Pub/Sub**, a shared ingestion service on **GKE**, orchestration across **Cloud Composer** and a legacy enterprise scheduler, modelling in **BigQuery** and **dbt**, the whole estate provisioned with **Terraform**.

The hard part is not the tools. It is making a modern cloud stack and a twenty-year-old on-premise estate agree with each other, reliably, every night — and writing the decisions down so the next team can see why.

THREE THINGS WORTH KNOWING

1 I have taken real money out of real systems.

A customer preference API I designed on GCP — App Engine, Pub/Sub, DataFlow, BigQuery — served millions of customers and **removed \$600,000 a year** in operating cost. I also migrated an enterprise data lake from **Teradata to BigQuery** and replaced licensed tooling with cloud-native across five stacks.

2 Award-winning payments work.

Five years at that bank, three platforms. Before the data work I put its card brands into **Apple Pay, Google Pay and Click to Pay**, then migrated the whole card-journey estate onto GCP microservices behind an API gateway. The self-serve fraud journey I architected won at the **Banking Tech Awards 2024** and the **Card & Payments Awards 2025** — it took fraud reporting out of the call-centre queue and gave it to the customer.

3 I ship applied AI in the open.

hoocode, a deterministic terminal coding agent **published to npm**, with a unified LLM API across 25+ providers. Underneath it, Rust: **embeddingsearchtools** — ONNX inference, an **HNSW index written from scratch** and BM25 fusion for hybrid retrieval — plus offline speech recognition and token-efficient web retrieval. github.com/kolisachint

THE STACK

GCP (BigQuery, Cloud Composer, GKE, Pub/Sub, DataFlow, Apigee) · Confluent Kafka · Apache Airflow · dbt · Terraform · Python, SQL · AWS · microservices and API architecture · Teradata, Oracle, Greenplum

WHAT I AM LOOKING FOR

A **Solution / Data / Cloud Architect** role where the **data platform is the product** rather than the delivery vehicle — product companies, payments businesses, GCCs. Pune or Bengaluru, or remote. Especially interested in teams putting ML or retrieval on top of a serious data platform — that is the layer I already build for myself.

BACKGROUND

18 years, all in enterprise data · **Google Cloud Certified Professional Data Engineer** · BE Computer Science & Engineering, Walchand College of Engineering, Sangli · Currently **Tata Consultancy Services**; previously Sears, Cognizant, Mahindra Satyam · Clients: **Lloyds Banking Group, Sears Holdings, Saudi Telecom, Barclays, Telstra**