

Sachin Koli

Solution Architect — Cloud Data & AI Platforms

Google Cloud Certified — Professional Data Engineer

Pune, India · +91 95522 36200 · kolisachint@gmail.com · linkedin.com/in/kolisachint · kolisachint.github.io

SUMMARY

Google Cloud Certified Professional Data Engineer with **18 years** building enterprise data platforms, now architecting **real-time streaming and batch data infrastructure on Google Cloud** for a UK tier-1 retail bank — Confluent Kafka, Cloud Composer, dbt, BigQuery and Terraform, spanning legacy on-premise and cloud environments.

Previously architected **digital cards and payments** at the same bank, migrating card servicing to GCP. The self-serve fraud journey I designed won at the **Banking Tech Awards 2024** and the **Card & Payments Awards 2025**.

Across twelve prior years in retail, telecom and banking data: a **Teradata** → **BigQuery** lake migration, and a customer preference API on GCP that **cut \$600,000 a year** in operating cost. I work the full arc — platform selection, architecture, build, and the stakeholder conversation that gets it funded.

EXPERIENCE

Tata Consultancy Services

Sep 2021 – Present

Solution Architect · Pune, India

Client: Lloyds Banking Group, United Kingdom

Data Platform & Real-Time Analytics — Jul 2025 – Present

- Architected the bank's **end-to-end batch and real-time streaming data platform** on Google Cloud, supporting high-throughput analytics, operational reporting and machine learning workloads.
- Engineered high-availability, event-driven ingestion pipelines on **Confluent Kafka** with Python, delivering low-latency data for operational reporting and instant business insight.
- Architected hybrid batch orchestration across **Cloud Composer (Apache Airflow)** and **IBM Tivoli Workload Scheduler**, managing complex dependency graphs that span legacy on-premise and cloud environments.
- Designed optimised **BigQuery** warehouse schemas and modular **dbt** transformation pipelines with enforced data quality and lineage; built automated testing, validation and monitoring frameworks so downstream analytics fail loudly rather than silently.
- Standardised infrastructure with **Terraform** — Composer environments, BigQuery datasets, Kafka connectors and IAM roles — and integrated it into **CI/CD**, cutting configuration drift and deployment cycle time.
- Implemented BigQuery and GCP **cost optimisation and governance**, reducing compute overhead while improving query performance and data security.
- Partnered with leadership, security and business units to turn requirements into production-ready architectural blueprints.

Digital Cards Modernisation — 2022 – Jul 2025

- Led end-to-end modernisation of **Digital Cards** — Add to Wallet, card controls, View PIN / View Card — migrating from on-premise infrastructure to **GCP**.
- Architected secure, scalable **microservices and cross-cloud API patterns** meeting banking compliance and performance requirements.
- Award-winning self-serve fraud journey** — recognised at the Banking Tech Awards 2024 and the Card & Payments Awards 2025 — moving fraud handling from an assisted call-centre process to a journey the customer completes.
- Designed service discovery, API gateways, containerisation, load balancing and secure inter-service communication across a hybrid estate.

Sears IT and Management Services India

Feb 2018 – Sep 2021

Technical Specialist → Architect · Pune, India

Client: Sears Holdings Corporation, United States — Targeted Interactions / TI Mart

- Built the **Email Preference Center API** on GCP App Engine, Pub/Sub, DataFlow and BigQuery, serving millions of customers — **reduced annual cost by \$600,000**.
- Migrated DataMarts from Teradata to BigQuery**, on schedule and without significant incident.
- Built an **API management layer** decoupling applications from cloud complexity across a multicloud estate — reused as the blueprint by other teams.
- Replaced self-hosted tooling with cloud-native across five stacks: Teradata → BigQuery · Unix → Compute Engine · Control-M → Airflow · shell/ETL → Python · GitLab Enterprise → Cloud Source Repositories.
- Designed shared datamarts (CDW, NSegment, Campaign) in BigQuery; led enterprise migration from **IBM Unica to Adobe Campaign** and Telluride, coordinating vendors (Adobe, BounceX) and internal teams.
- Delivered Tableau and Data Studio reporting; ran continuous GCP cost optimisation.

Cognizant Technology Solutions

Nov 2010 – Feb 2018

Senior Associate – Projects · Pune, India

Client: Saudi Telecom Company, Saudi Arabia — STC EDW (CLDM7 / CLDM12) · Onsite IT Delivery Lead · Feb 2015 – Feb 2018

- Led data modelling for the enterprise warehouse covering customers, accounts, subscriptions, offers and products across all channels and business units.
- Extended the model by **snowflaking dimensions** — demographics, account status, subscription status.
- Built an **automation tool** propagating a new key through every dimension table, Teradata utility script (FastLoad, MultiLoad, FastExport, BTEQ) and Ab-Initio object automatically.
- Ran a big-data POC on **Apache Hive and HDFS** for CDR ingestion.

Client: Barclays, United Kingdom — Emerging Markets DW & One Africa Program · Team Lead · Nov 2010 – Feb 2015

- Built staging from core banking (FLEX), ATM/POS (SPARROW) and online transaction sources, then the DataMarts above them.
- **One Africa:** unified emerging-market data across African countries into a single source of truth, letting one DataMart span new countries with no rework.
- Designed the **logical and physical data models**; built reusable automation — TWS object-creation tool, Teradata acquisition code generator, Unix deployment tooling.
- Implemented archiving and purging against the active warehouse, tiering history to **Greenplum**.

Mahindra Satyam

Jan 2008 – Nov 2010

Software Developer – S2 · Hyderabad / Bangalore, India

Client: Telstra Corporation, Australia — TR2.1 / TR2.2 EDW Data-In

- Built **DataStage 7.5** jobs, Teradata utility scripts and Perl/Unix tooling for the retail customer warehouse — registrations, demographics, segmentation, billed and unbilled call records.
- Built **Informatica** mappings and workflows migrating Oracle to Teradata for an enterprise data warehouse.

INDEPENDENT ENGINEERING — APPLIED AI

- **Rust embedding library** — a from-scratch implementation of text embedding generation, written to understand the mechanics rather than call an API.
- **Semantic code search** — retrieval over a codebase using vector similarity rather than keyword matching.

The data platforms I design are the natural substrate for retrieval and ML workloads; this is the layer I am moving toward.

TECHNICAL

Cloud	Google Cloud Platform — BigQuery, Cloud Composer, DataFlow, Pub/Sub, App Engine, Cloud Functions, DataProc, Cloud Storage, Compute Engine, Memorystore, Cloud SQL, VPC, IAM · AWS (S3)
Streaming	Confluent Kafka, Pub/Sub, Apache Beam, event-driven architecture
Orchestration	Cloud Composer / Apache Airflow, IBM Tivoli Workload Scheduler, Control-M
Transformation	dbt, Python, SQL, Apache Spark, Informatica, IBM DataStage, Ab-Initio
Infrastructure	Terraform, Git, CI/CD, containerisation
Architecture	Microservices, API gateways, service discovery, load balancing, multicloud and hybrid patterns, data governance, cost optimisation
Databases	BigQuery, Teradata, Oracle, Greenplum
BI	Tableau, Looker Studio, Cognos
Domains	Banking, cards & payments · Retail · Telecom

EDUCATION & CERTIFICATIONS

BE, Computer Science and Engineering — Walchand College of Engineering, Sangli (Shivaji University), 2003–2007

- **Google Cloud Certified — Professional Data Engineer**
- Google Cloud Professional Cloud Architect — *in progress*
- Data Engineering, Big Data and Machine Learning on GCP — Coursera specialisation
- Teradata Basics V2R5 · Teradata SQL V2R5 · Introduction to Oracle9i: SQL